

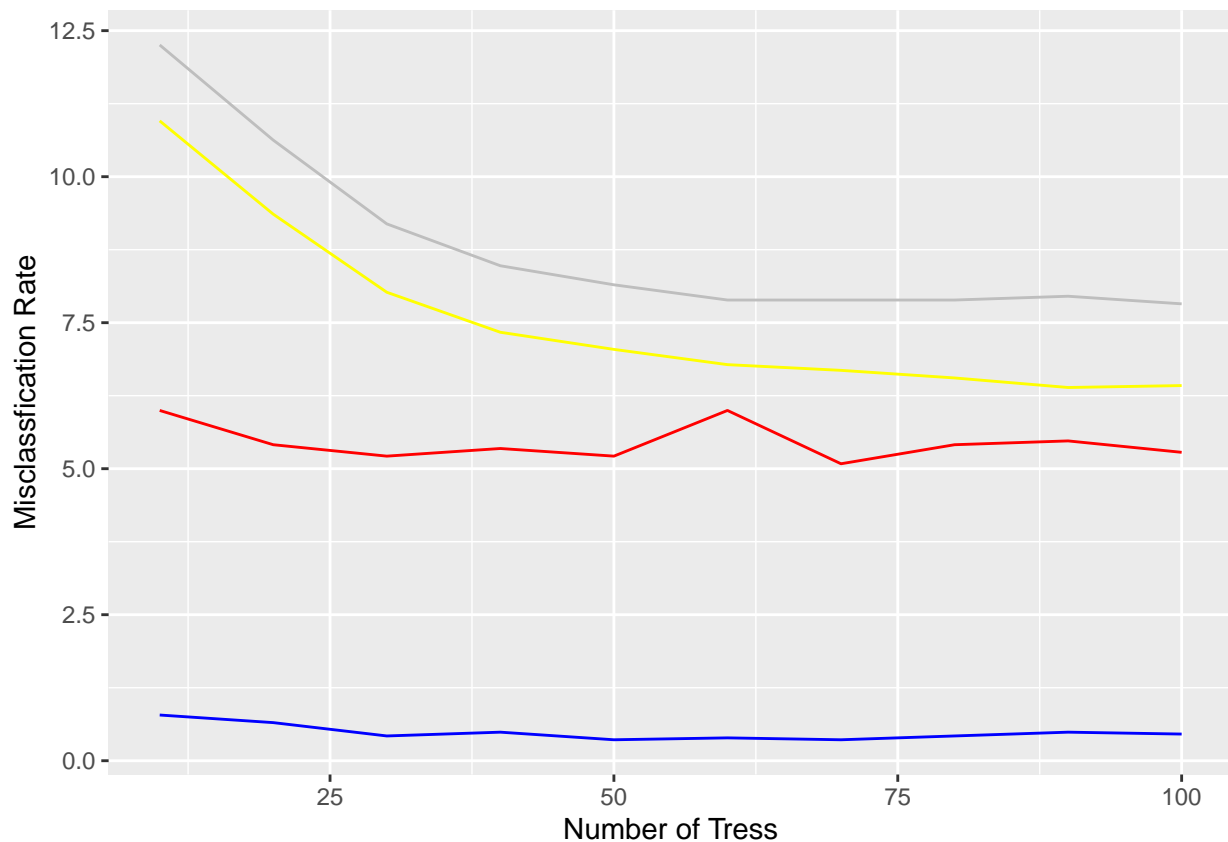
Block 2 Lab 1 Report

Zuxiang Li(zuxli371)

12/3/2019

1. ENSEMBLE METHODS

The file spambase.csv contains information about the frequency of various words, characters, etc. for a total of 4601 e-mails. Furthermore, these e-mails have been classified as spams (spam = 1) or regular e-mails (spam = 0). Your task is to evaluate the performance of Adaboost classification trees and random forests on the spam data. Specifically, provide a plot showing the error rates when the number of trees considered are 10, 20, . . . , 100. To estimate the error rates, use 2/3 of the data for training and 1/3 as hold-out test data. To learn Adaboost classification trees, use the function `blackboost()` of the R package `mboost`. Specify the loss function corresponding to Adaboost with the parameter family. To learn random forests, use the function `randomForest` of the R package `randomForest`.



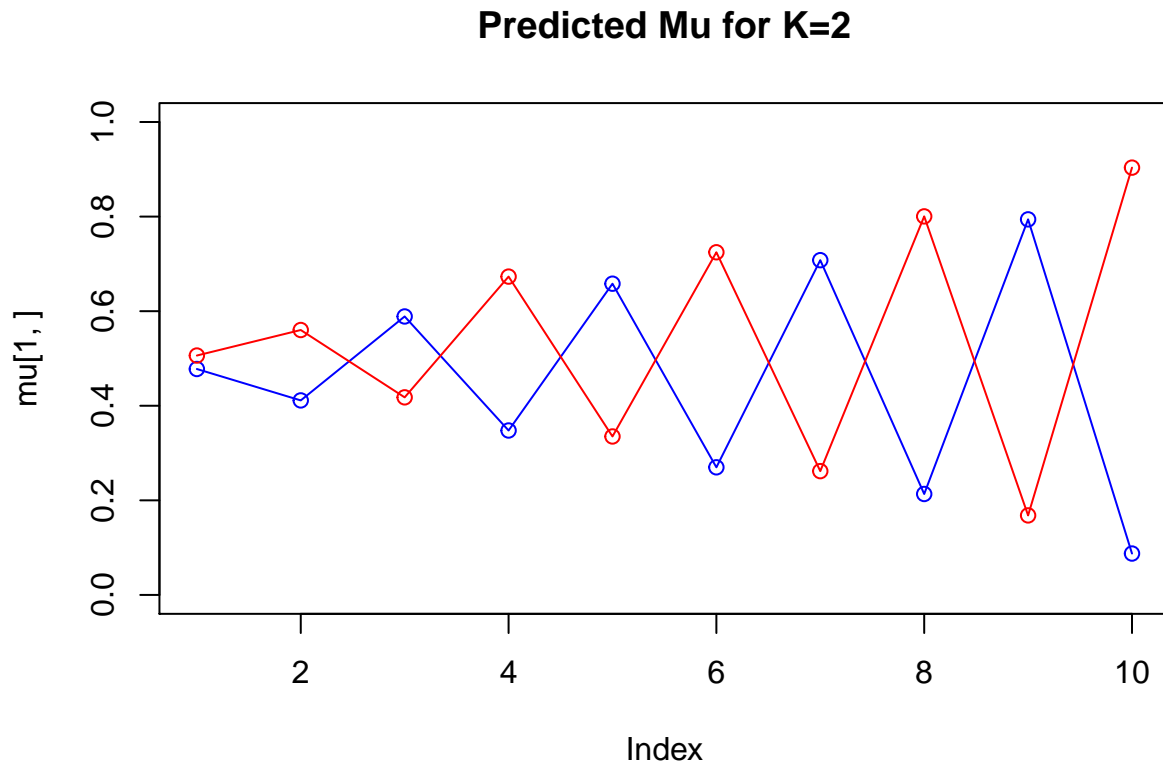
The plot above showed the Misclassification rate using different models and number of trees. Grey and yellow represent Adaboost model for train and test data. Red and blue lines are Randomforest model for train and test data.

The misclassification values of Adaboost with increasing number of trees are constantly decreasing for both train and test data. In every iteration a classifier is produced and given different weight. Then use the classifier generated in previous iteration to produce. It returns a weighted average of the classifiers. So the rate keeps decreasing.

As for Random forest, number of trees seems have no effect on the misclassification rate for both train and test data. Since Random forest used bagging, individual classifiers are then combined by taking a simple majority vote of their decisions. Random forest model have a better performance than Adaboost.

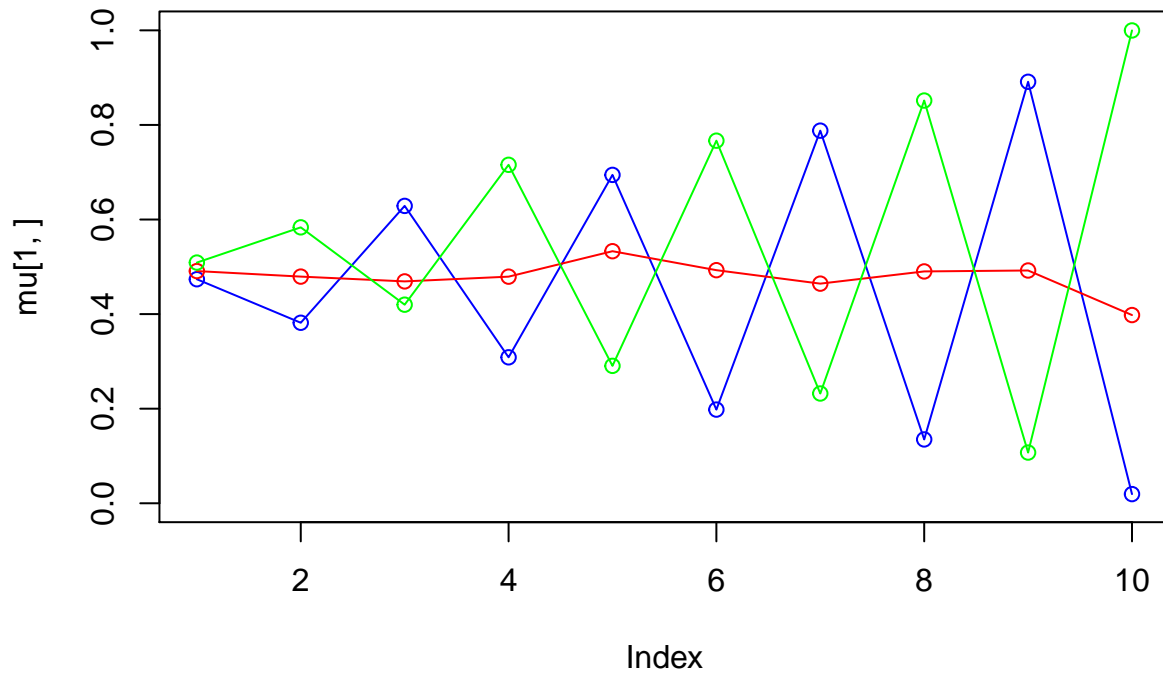
2. MIXTURE MODELS

Your task is to implement the EM algorithm for mixtures of multivariate Bernoulli distributions. Please use the template in the next page to solve the assignment. Then, use your implementation to show what happens when your mixture models has too few and too many components, i.e. set $K = 2, 3, 4$ and compare results. Please provide a short explanation as well.




```
## [1] 0.4981919 0.5018081
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4777136 0.4113065 0.5888317 0.3477062 0.6580979 0.2698303 0.7078467
## [2,] 0.5061835 0.5601552 0.4177868 0.6731171 0.3350702 0.7245255 0.2617664
##      [,8]      [,9]     [,10]
## [1,] 0.2134061 0.7941168 0.0875152
## [2,] 0.8004711 0.1681467 0.9035340
```

Predicted Mu for K=3




```
## [1] 0.3259592 0.3044579 0.3695828
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4737193 0.3817120 0.6288021 0.3086143 0.6943731 0.1980896 0.7879447
## [2,] 0.4909874 0.4793213 0.4691560 0.4791793 0.5329895 0.4928830 0.4643990
## [3,] 0.5089571 0.5834802 0.4199272 0.7157107 0.2905703 0.7667258 0.2320784
##      [,8]      [,9]     [,10]
## [1,] 0.1349651 0.8912534 0.01937869
## [2,] 0.4902682 0.4922194 0.39798407
## [3,] 0.8516111 0.1072226 0.99981353
```

Predicted Mu for K=4



```
## [1] 0.1614155 0.1383613 0.3609912 0.3392319
##      [,1]      [,2]      [,3]      [,4]      [,5]      [,6]      [,7]
## [1,] 0.4372908 0.5716691 0.6230114 0.4717152 0.4251232 0.2940734 0.4797605
## [2,] 0.5381955 0.3913346 0.2971686 0.5062848 0.6375272 0.7107583 0.4202372
## [3,] 0.5102441 0.5846281 0.4200464 0.7178717 0.2850900 0.7735833 0.2327656
## [4,] 0.4797762 0.3788928 0.6181216 0.3114748 0.6964392 0.2149967 0.7793732
##      [,8]      [,9]      [,10]
## [1,] 0.4812185 0.5945364 0.500815206
## [2,] 0.5246082 0.3534161 0.384513179
## [3,] 0.8546627 0.1022323 0.999999734
## [4,] 0.1450708 0.8791286 0.005800712
```



In this case, we generate data using `true_mu` and `true_pi`. For $K=2$, the blue dots, which means too few components, the blue points in the plot, only 16 iterations were made, this model is underfitted. For $K=4$, the grey dots, too many components, grey dots in the plot, 66 iterations were made, then this model is overfitted. Only when $K=3$, the red dots in plot, we got similar μ and π compare to the true values.

Appendix

1. ENSEMBLE METHODS

```
library(mboost)
library(randomForest)
# Import data
sp <- read.csv2("material/spambase.csv")
sp$Spam <- as.factor(sp$Spam)
n <- dim(sp)[1]
set.seed(1234567890)
id <- sample(1:n, floor(n*2/3))
train <- sp[id,]
test <- sp[-id,]

# Generate confusion matrix and calculate misclassification rate
cft<-function(Pred,Actual){
  cft<-table(Pred,Actual)
  # True positive
  tp <- cft[2, 2]
  # True negative
  tn <- cft[1, 1]
  # False positive
  fp <- cft[2, 1]
  # False negative
  fn <- cft[1, 2]

  misclassification <- (1-(tp + tn)/(tp + tn + fp + fn))*100
  return(misclassification)
}

ensemble_methods<-function(method,data){
  err_rate<-c()
  # select methods mboost
  if(substitute(method)=="mboost"){
    # Loop for 10 times each time the number of trees increase 10
    for(i in seq(10,100,10)){
      # Generate mboost model using train data,specify loss function as Binomial()
      bb<-blackboost(Spam~.,data=train,control = boost_control(mstop = i),family = AdaExp())
      # Get the predict value based on model we generated
      pred<-predict.mboost(bb,data,type="class")
      # Get the missclassification rate for current number of trees
      res<-cft(Pred=pred,Actual=data$Spam)
      err_rate<-c(err_rate,res)
    }
  }
  # select methods randomforest
  else if(substitute(method)=="randomforest"){
    # Loop for 10 times each time the number of trees increase 10
    for(i in seq(10,100,10)){
      # Generate randomforest model
```

```

    rf<-randomForest(Spam~.,data=train,ntree=i)
    # Get predict value
    pred<-predict(rf,data)
    # Get the missclassification rate for current number of trees
    res<-cft(Pred=pred,Actual=data$Spam)
    err_rate<-c(err_rate,res)
  }
}
return(err_rate)
}
err_rate1<-ensemble_methods("mboost",train)
err_rate2<-ensemble_methods("mboost",test)
err_rate3<-ensemble_methods("randomforest",train)
err_rate4<-ensemble_methods("randomforest",test)

df<-data.frame(step=seq(10,100,10),err1=err_rate1,err2=err_rate2,err3=err_rate3,err4=err_rate4)
ggplot()+xlab("")+geom_line(data = df, aes(x = step, y = err3), color = "blue") +
  geom_line(data = df, aes(x = step, y = err4), color = "red") +geom_line(data = df, aes(x = step, y = err1), color = "grey")
  geom_line(data = df, aes(x = step, y = err2), color = "grey")+xlab("Number of Tress")+ylab("Misclassification Rate")

```

2. MIXTURE MODELS

```
RNGversion('3.5.1')
```

```
## Warning in RNGkind("Mersenne-Twister", "Inversion", "Rounding"): non-uniform
## 'Rounding' sampler used
```

```

EM_algorithm<-function(K){
  set.seed(1234567890)
  max_it <- 100 # max number of EM iterations
  min_change <- 0.1 # min change in log likelihood between two consecutive EM iterations
  N=1000 # number of training points
  D=10 # number of dimensions
  x <- matrix(nrow=N, ncol=D) # training data
  true_pi <- vector(length = 3) # true mixing coefficients
  true_mu <- matrix(nrow=3, ncol=D) # true conditional distributions
  true_pi=c(1/3, 1/3, 1/3)
  true_mu[1,]=c(0.5,0.6,0.4,0.7,0.3,0.8,0.2,0.9,0.1,1)
  true_mu[2,]=c(0.5,0.4,0.6,0.3,0.7,0.2,0.8,0.1,0.9,0)
  true_mu[3,]=c(0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5,0.5)
  plot(true_mu[1,], type="o", col="blue", ylim=c(0,1))
  points(true_mu[2,], type="o", col="red")
  points(true_mu[3,], type="o", col="green")
  # Producing the training data
  for(n in 1:N) {
    k <- sample(1:3,1,prob=true_pi)
    for(d in 1:D) {
      x[n,d] <- rbinom(1,1,true_mu[k,d])
    }
  }
  #K=3 # number of guessed components
}

```

```

z <- matrix(nrow=N, ncol=K) # fractional component assignments
pi <- vector(length = K) # mixing coefficients
mu <- matrix(nrow=K, ncol=D) # conditional distributions
llik <- vector(length = max_it) # log likelihood of the EM iterations
# Random initialization of the parameters
pi <- runif(K,0.49,0.51)
pi <- pi / sum(pi)
for(k in 1:K) {
  mu[k,] <- runif(D,0.49,0.51)
}

for(it in 1:max_it) {
  plot(mu[1,], type="o", col="blue", ylim=c(0,1))
  points(mu[2,], type="o", col="red")
  points(mu[3,], type="o", col="green")
  #points(mu[4,], type="o", col="yellow")
  Sys.sleep(0.5)
  # E-step: Computation of the fractional component assignments

  for(n in 1:N){
    denominator<-c()
    numerator<-c()
    for(k in 1:K){
      new_numerator<-pi[k]*prod(mu[k,]^x[n,]*(1-mu[k,])^(1-x[n,]))
      numerator<-c(numerator,new_numerator)
      denominator<-sum(numerator)
    }
    z[n,]<-numerator/denominator
  }
  #Log likelihood computation.
  sum=0
  for(n in 1:N){
    for(k in 1:K){
      sum=sum+z[n,k]*(log(pi[k])+sum(x[n,]*log(mu[k,])+(1-x[n,])*log(1-mu[k,])))
    }
  }
  llik[it]<-sum

  cat("iteration: ", it, "log likelihood: ", llik[it], "\n")
  flush.console()
  # Stop if the log likelihood has not changed significantly
  if(it>1){
    if(abs(llik[it]-llik[it-1])<min_change)
      break
  }

  #M-step: ML parameter estimation from the data and fractional component assignments
  # New mixing coefficients for next iteration
  pi=apply(z,2,mean)
  # New conditional distributions
  for(k in 1:K){
    sum=0
    for (n in 1:N) {

```





```

        sum = sum + x[n,] * z[n,k]
    }
    mu[k,] = sum / sum(z[,k])
}
}
pi
mu
plot(llik[1:it], type="o")
return(llik[1:it])
#return(list(pi=pi,mu=mu,res=,llik=llik[1:it]))
}


#em1<-EM_algorithm(K=2)
em2<-EM_algorithm(K=3)

```







iteration: 1 log likelihood: -8029.723




iteration: 2 log likelihood: -8027.183




iteration: 3 log likelihood: -8024.696




iteration: 4 log likelihood: -8005.631




iteration: 5 log likelihood: -7877.606




iteration: 6 log likelihood: -7403.513




iteration: 7 log likelihood: -6936.919




iteration: 8 log likelihood: -6818.582




iteration: 9 log likelihood: -6791.377




iteration: 10 log likelihood: -6780.713




iteration: 11 log likelihood: -6774.958




iteration: 12 log likelihood: -6771.261




iteration: 13 log likelihood: -6768.606




iteration: 14 log likelihood: -6766.535




iteration: 15 log likelihood: -6764.815




iteration: 16 log likelihood: -6763.316




iteration: 17 log likelihood: -6761.967




iteration: 18 log likelihood: -6760.727




iteration: 19 log likelihood: -6759.572




iteration: 20 log likelihood: -6758.491




iteration: 21 log likelihood: -6757.475




iteration: 22 log likelihood: -6756.521




iteration: 23 log likelihood: -6755.625




iteration: 24 log likelihood: -6754.784




iteration: 25 log likelihood: -6753.996




iteration: 26 log likelihood: -6753.26




iteration: 27 log likelihood: -6752.571




iteration: 28 log likelihood: -6751.928




iteration: 29 log likelihood: -6751.328




iteration: 30 log likelihood: -6750.768




iteration: 31 log likelihood: -6750.246




iteration: 32 log likelihood: -6749.758




iteration: 33 log likelihood: -6749.304




iteration: 34 log likelihood: -6748.88




iteration: 35 log likelihood: -6748.484




iteration: 36 log likelihood: -6748.114




iteration: 37 log likelihood: -6747.767




iteration: 38 log likelihood: -6747.444




iteration: 39 log likelihood: -6747.14




iteration: 40 log likelihood: -6746.856




iteration: 41 log likelihood: -6746.589




iteration: 42 log likelihood: -6746.338




iteration: 43 log likelihood: -6746.102




iteration: 44 log likelihood: -6745.88




iteration: 45 log likelihood: -6745.67




iteration: 46 log likelihood: -6745.472




iteration: 47 log likelihood: -6745.285




iteration: 48 log likelihood: -6745.108




iteration: 49 log likelihood: -6744.939




iteration: 50 log likelihood: -6744.78




iteration: 51 log likelihood: -6744.627




iteration: 52 log likelihood: -6744.483




iteration: 53 log likelihood: -6744.344




iteration: 54 log likelihood: -6744.212




iteration: 55 log likelihood: -6744.086




iteration: 56 log likelihood: -6743.964




iteration: 57 log likelihood: -6743.848




iteration: 58 log likelihood: -6743.736




iteration: 59 log likelihood: -6743.628




iteration: 60 log likelihood: -6743.524



iteration: 61 log likelihood: -6743.423



iteration: 62 log likelihood: -6743.326



```
#em3<-EM_algorithm(K=4)

# ggplot()+geom_point(aes(x=c(1:length(em1)),y=em1),color="blue")+
#   geom_point(aes(x=c(1:length(em2)),y=em2),color="red")+
#   geom_point(aes(x=c(1:length(em3)),y=em3),color="grey")+
#   xlab("Number of Iterations")+ylab("Expected Log Likelihood")
```